

# Exercise #10

## Principal Component Analysis

Brian Schweigler; 16-102-071

24/05/2022

### Preliminaries

Set a seed for later:

```
set.seed(1786397)
```

Loading the Boston dataset, set CHAS as a factor and show an overview:

```
setwd(dirname(rstudioapi::getActiveDocumentContext()$path))
boston_df = read.csv("Boston.txt", sep = " ", header = TRUE)
boston_df$chas <-
  factor(
    boston_df$chas,
    levels = c("0", "1"),
    labels = c("0", "1")
  )
summary(boston_df)
```

```
##      crim              zn      indus      chas      nox
## Min.   : 0.00632   Min.   : 0.00   Min.   : 0.46   0:471   Min.   :0.3850
## 1st Qu.: 0.08205   1st Qu.: 0.00   1st Qu.: 5.19   1: 35   1st Qu.:0.4490
## Median : 0.25651   Median : 0.00   Median : 9.69           Median :0.5380
## Mean   : 3.61352   Mean   : 11.36   Mean   :11.14           Mean   :0.5547
## 3rd Qu.: 3.67708   3rd Qu.: 12.50   3rd Qu.:18.10           3rd Qu.:0.6240
## Max.   :88.97620   Max.   :100.00   Max.   :27.74           Max.   :0.8710
##      rm      age      dis      rad
## Min.   :3.561   Min.   : 2.90   Min.   : 1.130   Min.   : 1.000
## 1st Qu.:5.886   1st Qu.: 45.02   1st Qu.: 2.100   1st Qu.: 4.000
## Median :6.208   Median : 77.50   Median : 3.207   Median : 5.000
## Mean   :6.285   Mean   : 68.57   Mean   : 3.795   Mean   : 9.549
## 3rd Qu.:6.623   3rd Qu.: 94.08   3rd Qu.: 5.188   3rd Qu.:24.000
## Max.   :8.780   Max.   :100.00   Max.   :12.127   Max.   :24.000
##      tax      ptratio      lstat      medv
## Min.   :187.0   Min.   :12.60   Min.   : 1.73   Min.   : 5.00
## 1st Qu.:279.0   1st Qu.:17.40   1st Qu.: 6.95   1st Qu.:17.02
## Median :330.0   Median :19.05   Median :11.36   Median :21.20
## Mean   :408.2   Mean   :18.46   Mean   :12.65   Mean   :22.53
## 3rd Qu.:666.0   3rd Qu.:20.20   3rd Qu.:16.95   3rd Qu.:25.00
## Max.   :711.0   Max.   :22.00   Max.   :37.97   Max.   :50.00
```

No NAs; the range of the values can't be gauged without further information.

R somehow already removed the IDs, thus we will not need to do this.

## 1 Normalize your dataset and consider all the variables except MEDV. Create a PCA model and plot it.

First we will normalize the dataset:

```
normalize <- function(x) {  
  mean_x = mean(x)  
  return ((x - mean_x) / sd(x))  
}  
  
boston_df_numeric = lapply(boston_df, as.numeric)  
boston_normalized = as.data.frame(lapply(boston_df_numeric[1:13], normalize))  
boston_normalized$medv = boston_df$medv  
summary(boston_normalized)
```

```
##      crim      zn      indus      chas  
## Min.   :-0.419367 Min.   :-0.48724 Min.   :-1.5563 Min.   :-0.2723  
## 1st Qu.: -0.410563 1st Qu.: -0.48724 1st Qu.: -0.8668 1st Qu.: -0.2723  
## Median :-0.390280 Median :-0.48724 Median :-0.2109 Median :-0.2723  
## Mean   : 0.000000 Mean   : 0.00000 Mean   : 0.0000 Mean   : 0.0000  
## 3rd Qu.: 0.007389 3rd Qu.: 0.04872 3rd Qu.: 1.0150 3rd Qu.: -0.2723  
## Max.   : 9.924110 Max.   : 3.80047 Max.   : 2.4202 Max.   : 3.6648  
##      nox      rm      age      dis  
## Min.   :-1.4644 Min.   :-3.8764 Min.   :-2.3331 Min.   :-1.2658  
## 1st Qu.: -0.9121 1st Qu.: -0.5681 1st Qu.: -0.8366 1st Qu.: -0.8049  
## Median :-0.1441 Median :-0.1084 Median : 0.3171 Median :-0.2790  
## Mean   : 0.0000 Mean   : 0.0000 Mean   : 0.0000 Mean   : 0.0000  
## 3rd Qu.: 0.5981 3rd Qu.: 0.4823 3rd Qu.: 0.9059 3rd Qu.: 0.6617  
## Max.   : 2.7296 Max.   : 3.5515 Max.   : 1.1164 Max.   : 3.9566  
##      rad      tax      ptratio      lstat  
## Min.   :-0.9819 Min.   :-1.3127 Min.   :-2.7047 Min.   :-1.5296  
## 1st Qu.: -0.6373 1st Qu.: -0.7668 1st Qu.: -0.4876 1st Qu.: -0.7986  
## Median :-0.5225 Median :-0.4642 Median : 0.2746 Median :-0.1811  
## Mean   : 0.0000 Mean   : 0.0000 Mean   : 0.0000 Mean   : 0.0000  
## 3rd Qu.: 1.6596 3rd Qu.: 1.5294 3rd Qu.: 0.8058 3rd Qu.: 0.6024  
## Max.   : 1.6596 Max.   : 1.7964 Max.   : 1.6372 Max.   : 3.5453  
##      medv  
## Min.   : 5.00  
## 1st Qu.:17.02  
## Median :21.20  
## Mean   :22.53  
## 3rd Qu.:25.00  
## Max.   :50.00
```

Now we can calculate the PCA:

```
boston_pca <- princomp(boston_normalized[, -ncol(boston_normalized)], cor = TRUE)  
summary(boston_pca, loadings = T)
```

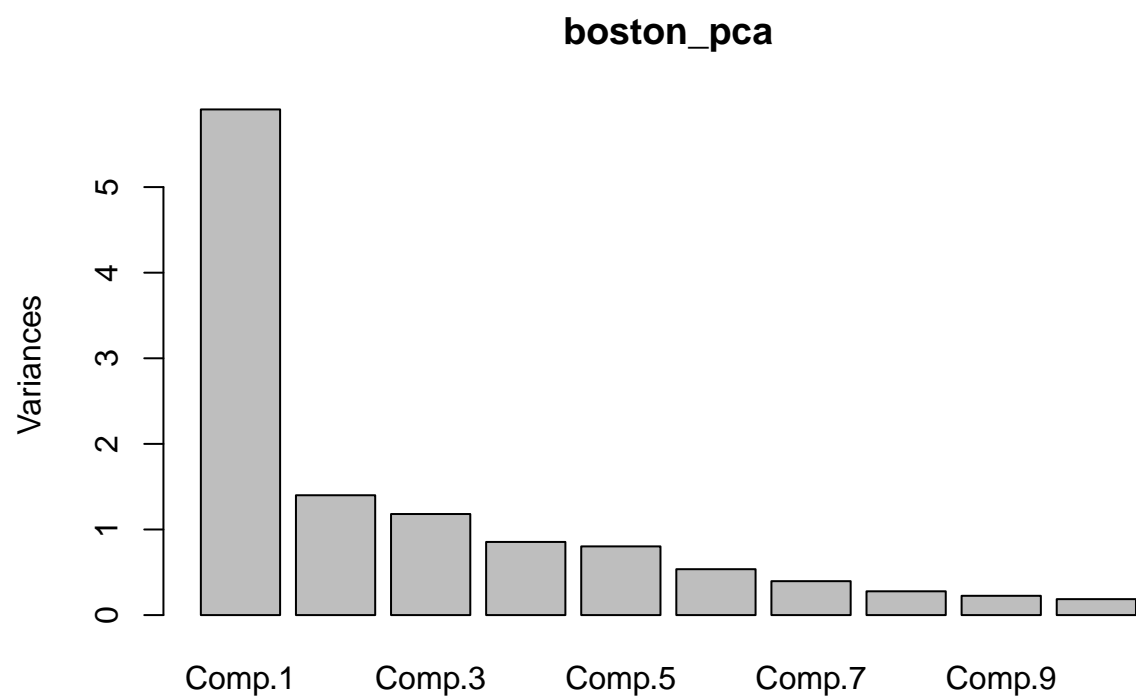
```
## Importance of components:  
##      Comp.1  Comp.2  Comp.3  Comp.4  Comp.5  
## Standard deviation 2.4304497 1.1832423 1.08659863 0.92436016 0.8957392  
## Proportion of Variance 0.4922571 0.1166719 0.09839138 0.07120348 0.0668624  
## Cumulative Proportion 0.4922571 0.6089290 0.70732039 0.77852386 0.8453863  
##      Comp.6  Comp.7  Comp.8  Comp.9  Comp.10  
## Standard deviation 0.7322458 0.6293692 0.52723411 0.47451400 0.43151563
```

```

## Proportion of Variance 0.0446820 0.0330088 0.02316465 0.01876363 0.01551715
## Cumulative Proportion 0.8900683 0.9230771 0.94624170 0.96500533 0.98052248
##                               Comp.11      Comp.12
## Standard deviation      0.41256242 0.252036760
## Proportion of Variance 0.01418398 0.005293544
## Cumulative Proportion 0.99470646 1.000000000
##
## Loadings:
##      Comp.1 Comp.2 Comp.3 Comp.4 Comp.5 Comp.6 Comp.7 Comp.8 Comp.9 Comp.10
## crim      0.251 0.274 0.351      0.192 0.760 0.155 0.272      0.109
## zn       -0.266 0.250 0.359 0.194 0.402 -0.295 -0.401 0.374 0.259 -0.268
## indus     0.355      -0.503 0.201 0.811 -0.197      -0.345 0.174 0.633 -0.374 0.313
## chas      -0.503 0.201 0.811 -0.197
## nox       0.350 -0.233      0.215 -0.208      0.203 -0.136
## rm       -0.196 -0.273 0.561 -0.402 -0.285      -0.327      -0.441
## age       0.323 -0.293      -0.163      0.109 -0.600      0.392 0.456
## dis      -0.331 0.343      0.234      -0.122 -0.162 -0.166 0.690
## rad       0.322 0.231 0.408      -0.119 -0.149      -0.462
## tax       0.342 0.213 0.331      -0.329      -0.168      0.115
## ptratio  0.211 0.393 -0.184 0.109 -0.702      -0.318 0.255 0.126 -0.186
## lstat     0.315 0.128 -0.264 0.185 0.346      -0.425 -0.220 -0.598 -0.255
##      Comp.11 Comp.12
## crim
## zn
## indus    -0.116 -0.251
## chas
## nox      0.807
## rm       0.135
## age     -0.188
## dis      0.402
## rad     -0.115 -0.633
## tax     -0.221 0.721
## ptratio 0.214
## lstat

```

plot(boston\_pca)



As should be the case, component 1 has the highest impact. It is interesting how not all variables impact all the different components.

## 2. Which predictor variable contributes the most to component 1? And which contributes the least?

This is quite straightforward:

```
comp1_max_contribution = max(abs(boston_pca$loadings[,1]))
comp1_max_contribution
```

```
## [1] 0.3548731
```

```
max_row_name = which(abs(boston_pca$loadings[,1]) == comp1_max_contribution)
max_row_name
```

```
## indus
```

```
##      3
```

So indus has the highest impact on PCA's component 1.

```
comp1_min_contribution = min(abs(boston_pca$loadings[,1]))
comp1_min_contribution
```

```
## [1] 0.007588523
```

```
min.row.num = which(abs(boston_pca$loadings[,1]) == comp1_min_contribution)
min.row.num
```

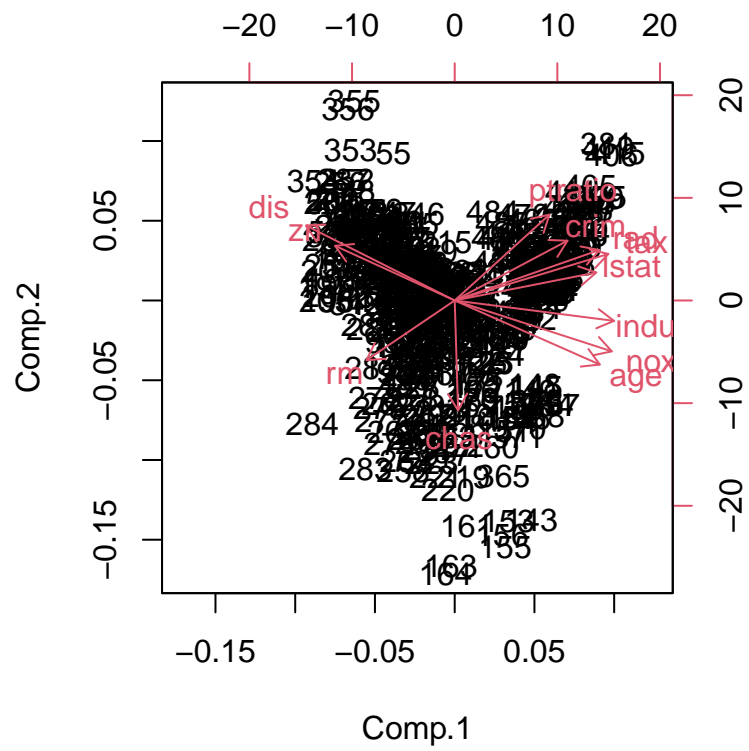
```
## chas
```

```
##      4
```

The smallest contribution to component 1 is from chas, but I am unsure if this is due to the factorizing or if it worked properly.

With a biplot, we also see that indus has barely the largest contribution to comp 1 and chas the smallest. But of note is that CHAS has the largest contribution to component 2.

```
biplot(boston_pca)
```



3. Estimate the proportion of variance explained by all the components. If we want to explain only 80% of the original data, how many components should we use?

```
summary(boston_pca)
```

```
## Importance of components:
##               Comp.1   Comp.2   Comp.3   Comp.4   Comp.5
## Standard deviation  2.4304497 1.1832423 1.08659863 0.92436016 0.8957392
## Proportion of Variance 0.4922571 0.1166719 0.09839138 0.07120348 0.0668624
## Cumulative Proportion 0.4922571 0.6089290 0.70732039 0.77852386 0.8453863
##               Comp.6   Comp.7   Comp.8   Comp.9   Comp.10
## Standard deviation  0.7322458 0.6293692 0.52723411 0.47451400 0.43151563
## Proportion of Variance 0.0446820 0.0330088 0.02316465 0.01876363 0.01551715
## Cumulative Proportion 0.8900683 0.9230771 0.94624170 0.96500533 0.98052248
##               Comp.11   Comp.12
## Standard deviation  0.41256242 0.252036760
## Proportion of Variance 0.01418398 0.005293544
## Cumulative Proportion 0.99470646 1.000000000
```

We need 5 of the components to boserve at least 80% of the variance in the data.

In general, the standard deviation and proportional variation can be also seen as follows:

```
boston_pca$sdev
```

```
##   Comp.1   Comp.2   Comp.3   Comp.4   Comp.5   Comp.6   Comp.7   Comp.8
## 2.4304497 1.1832423 1.0865986 0.9243602 0.8957392 0.7322458 0.6293692 0.5272341
##   Comp.9   Comp.10   Comp.11   Comp.12
## 0.4745140 0.4315156 0.4125624 0.2520368
```

```
var.proportion = (boston_pca$sdev)^2/sum((boston_pca$sdev)^2)
var.proportion
```

```
##   Comp.1   Comp.2   Comp.3   Comp.4   Comp.5   Comp.6
## 0.492257148 0.116671856 0.098391382 0.071203475 0.066862395 0.044681998
##   Comp.7   Comp.8   Comp.9   Comp.10   Comp.11   Comp.12
## 0.033008799 0.023164650 0.018763628 0.015517145 0.014183979 0.005293544
```

#### 4. Generate a new dataset using only the components selected in Problem 3. Create a multiple regression model using these components as predictors for the target variable MEDV

Now we divide this into test (30%) and train (70%) data set:

```
# Generate the training and test sets
boston_pca_data = data.frame(boston_pca$scores, medv = boston_df$medv)
train_size = round(nrow(boston_pca_data)*0.7)
train_index = sample( c(1: nrow(boston_pca_data)), train_size)
boston_train = boston_pca_data[train_index,]
boston_test = boston_pca_data[-train_index,]
```

Now we will only use the first 5 components:

```
boston_small = boston_train[,c(1:5,13)]
pca_lm_small = lm(medv ~ ., data = boston_small)
summary(pca_lm_small)
```

```
##
## Call:
## lm(formula = medv ~ ., data = boston_small)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -18.581  -2.979  -0.806   1.883  33.401
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  22.5063     0.2941  76.527 < 2e-16 ***
## Comp.1       -2.2487     0.1237 -18.184 < 2e-16 ***
## Comp.2       -2.6336     0.2486 -10.593 < 2e-16 ***
## Comp.3        3.0910     0.2743  11.267 < 2e-16 ***
## Comp.4       -1.7467     0.3297  -5.298 2.08e-07 ***
## Comp.5       -1.3761     0.3333  -4.129 4.58e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 5.519 on 348 degrees of freedom
## Multiple R-squared:  0.6412, Adjusted R-squared:  0.6361
## F-statistic: 124.4 on 5 and 348 DF,  p-value: < 2.2e-16

MSE_train_lm_small = mean(pca_lm_small$residuals^2)
print(sprintf("MSE Train Small = %f", MSE_train_lm_small))

## [1] "MSE Train Small = 29.945893"

boston_small_prediction = predict(pca_lm_small, boston_test)
MSE_test_lm_small = mean(boston_small_prediction^2)
print(sprintf("MSE Test Small = %f", MSE_test_lm_small))

## [1] "MSE Test Small = 569.174042"
```

Test MSE is worse than Train MSE, so this is not a useful model.



## 5. Compare the model created in Problem 4 with a multiple regression model using all the components

Let's get to it:

```
pca_lm_full = lm(medv ~ ., data = boston_train)
summary(pca_lm_full)

##
## Call:
## lm(formula = medv ~ ., data = boston_train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -14.3370  -3.2142  -0.6896   2.2473  24.9651
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  22.44952    0.27171  82.622 < 2e-16 ***
## Comp.1       -2.29097    0.11485 -19.947 < 2e-16 ***
## Comp.2       -2.75957    0.23111 -11.941 < 2e-16 ***
## Comp.3        3.09703    0.25430  12.179 < 2e-16 ***
## Comp.4       -1.77364    0.30557  -5.804 1.48e-08 ***
## Comp.5       -1.53765    0.30836  -4.986 9.81e-07 ***
## Comp.6       -0.04797    0.39284  -0.122  0.90289
## Comp.7        1.10234    0.42830   2.574  0.01048 *
## Comp.8        0.06862    0.52522   0.131  0.89614
## Comp.9        1.76789    0.55495   3.186  0.00158 **
## Comp.10       -1.43454    0.63560  -2.257  0.02464 *
## Comp.11       -3.94167    0.66052  -5.968 6.04e-09 ***
## Comp.12       -3.54407    1.08787  -3.258  0.00124 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 5.084 on 341 degrees of freedom
## Multiple R-squared:  0.7017, Adjusted R-squared:  0.6912
## F-statistic: 66.84 on 12 and 341 DF, p-value: < 2.2e-16

MSE_train_lm_full = mean(pca_lm_full$residuals^2)
print(sprintf("MSE Train Full = %f", MSE_train_lm_full))

## [1] "MSE Train Full = 24.900369"

boston_full_prediction = predict(pca_lm_full, boston_test)
MSE_test_lm_full = mean(boston_full_prediction^2)
print(sprintf("MSE Test Full = %f", MSE_test_lm_full))

## [1] "MSE Test Full = 574.226813"
```

Welp, this model is also not all that useable. Slightly better performance than the small one, but still not all that useable in the real world.

Nonetheless, the full model does not perform all that much better than the small one, which should be noted.